


Graph Algorithms (Greedy)

Lecture 6

Timothy Kim
GWU CSCI 6212

Minimum Spanning Tree

- **Problem:** Connect all vertices together deeply as possible
- **Input:** undirected graph $G=(V,E)$ and c_e for $e \in E$
- **Output:**
 - Minimum Cost Tree $T \subseteq E$ that spans V
 - No cycle
 - Connected
- **Assumptions:**
 - G is connected to begin with
 - C is unique



Prim's Algorithm

```
1. Prim(G):  
2.   X = {s}  // s is chosen arbitrarily  
3.   T = {}  
4.   while X ≠ V:  
5.     let e = (u,v)  
6.     where e is the cheapest crossing edge of cut (X, V-X)  
7.     T = T + e  
8.     X = X + v  
9.   return T
```

Proof of Correctness

- **Claim:** Prim's algorithm correctly computes an MST
- **Part I:** Prim's algorithm produces a spanning tree T^*
 - Spanning = all vertices are included
 - Tree = no cycles
- **Part II:** T^* is a MST
 - Minimal cost

Definitions to recall

- **Connected Graph:** a graph there is a path between every pair of vertices

Part I


- **Empty cut lemma:**
 - Graph, G , is not connected $\iff \exists$ cut (A,B) of G with no crossing edges
- **Proof of \Leftarrow :**
 - Assume \exists cut (A,B) with no crossing edges
 - Pick any $u \in A$ and $v \in B$
 - Since there are no crossing edges given cut (A,B) , there is no edge (u,v)
 - \therefore by the definition of connected graph, G is not connected

Part I cont.

- **Empty cut lemma:**
 - Graph, G , is not connected $\Leftrightarrow \exists$ cut (A,B) of G with no crossing edges
- **Proof of \Rightarrow :**
 - Assume Graph, G , is not connected
 - Pick any $u \in G$
 - Create a cut of (A, B) such that
 - $A = \{ \text{all vertices reachable from } u \}$
 - $B = \{ \text{all other vertices} \}$
 - Then cut (A, B) has no crossing edges

Part 1 cont.

- **Double Crossing Lemma:**
 - Suppose cycle, $C \subseteq E$, has an edge crossing a cut (A,B) then there must exist another edge that crosses that cut.
 - Proof by contradiction
- **Lemma 3:**
 - If e is the only edge crossing a cut (A,B) then e is not part of any cycle.
 - Proof using Double Crossing Lemma



Part 1 cont.


- **Part I:** Prim's algorithm produces a **spanning** tree
 - The algorithm chooses only edges stemming from X . Therefore the algorithm maintains the invariant of T spans X (meaning T includes all the vertices in X).
 - Algorithm must halt (X eventually $= V$), otherwise the cut of $(X, V - X)$ will have no crossing edges. If the cut has no crossing edges, by Empty Cut Lemma, G must be disconnected. This is a contradiction, thus the algorithm halts with $X = V$.
- \therefore Prim's algorithm produces a T that spans V

Part I cont.

- **Part I:** Prim's algorithm produces a spanning **tree** (no cycle)
 - Whenever an edge, e , gets added to T , e is the first edge to cross the cut $(X, V - X)$. By Lemma 3, e does not create a cycle
 - \therefore Prim's algorithm produces a tree

Part II

- **Part II:** T^* is a MST
 - Minimal cost
- **The Cut Property**
 - Given $e \in G$, suppose \exists cut (A,B) | e is the cheapest crossing edge, then $e \in \text{MST}(G)$



Part II

- **Claim:** Cut Property \Rightarrow Prim's Algorithm produces MST(G)
- Every edge $e \in T^*$ is chosen as the cheapest crossing edge of cut $(X, V - X)$.
- By the cut property, $T^* \subseteq \text{MST}(G)$.
- From Part I, since T^* is a spanning tree of G , $T^* = \text{MST}(G)$. QED

Proof of cut property

- **The Cut Property**
 - Given $e \in G$, suppose \exists cut (A,B) | e is the cheapest crossing edge, then $e \in \text{MST}(G)$, Y
- By contradiction.
 - Suppose e is the cheapest crossing edge of a cut (A,B) of G , yet $e \notin \text{MST}(G)$, Y

Proof of cut property

- Suppose e is the cheapest crossing edge of a cut (A,B) of G , yet $e \notin \text{MST}(G)$, Y
- Since Y doesn't include e , then it must include another edge, $f \mid c_f > c_e$, $f \in Y$ and cuts (A,B) and is part of Y . (Otherwise, Y is not connected.)
- (We want to use a swap method here, but if e or f is part of a cycle, we can't just swap.)
- Since $f \in Y$, and Y is a spanning tree, $Y + e$ could create a cycle.

Proof of cut property


- By Double Cut Lemma, there must exist another edge, $e' \mid c_{e'} > c_e$ and creates the cycle with e
- (Now we can swap e with e')
- Note, $Y' = Y + e - e'$ is a spanning tree.
- But cost of $Y' < Y$ which is a contradiction. QED.

Prim's Algorithm Running Time

```
1. Prim(G):
2.    $X = \{s\}$  //  $s$  is chosen arbitrarily
3.    $T = \{\}$ 
4.   while  $X \neq V$ :
5.     let  $e = (u,v)$ 
6.     where  $e$  is the cheapest crossing edge of cut  $(X, V-X)$ 
7.      $T = T + e$ 
8.      $X = X + v$ 
9.   return  $T$ 
```


Running time of Prim's

- $O(n \times m)$ - Literal implementation
- We can use MinHeap where most of its operations are in $O(\log n)$
- Use Heap to store edges $\Rightarrow O(m \log n)$
- But faster to store vertices in Heap with following invariant
 - Invariant #1: Elements in heap = $v \in V - X$
 - Invariant #2: for $v \in V - X$:
 $\text{key}_v = \text{cheapest edge } (u, v)$




Prim's with MinHeap of Vertices

- **Preprocessing** (Initialization) of heap:
 - Initial cut = $(\{s\}, V - \{s\})$
 - Find all the edges that cross that cut and create heap:
 - $O(m + n \log n)$ or $O(m + n)$ if you use heapify
 - Since $m \geq n - 1$, $O(m + n) = O(n)$

Prim's with MinHeap of Vertices


- **During execution:**
 - To pick the cheapest edge, single ExtractMin call to heap will give you the right edge.
 - $O(\log n)$
- **Keeping the invariant**
after ExtractMin has been called



Prim's with MinHeap of Vertices

- **Keeping the invariant** after ExtractMin has been called
- Use extra metadata to speed up delete

1. When v is added to X :
2. for each edge $(v, w) \in E$:
3. if $w \in V - X$:
4. delete w from heap
5. $\text{key}[w] = \min\{\text{key}[w], c_{v,w}\}$
6. insert w into heap



Final Running Time of Prim






- Preprocessing: $O(n)$
- One ExtractMin called per each vertex: $O(n \log n)$
- Each edge triggers at most one delete/insert: $O(m \log n)$
- Entire running time = $O(m \log n) = O(n \log n)$
 - At the cost of factor of $\log n$, we can calculate MST as fast as reading in the graph.

Kruskal's MST

```
1. Kruskal(E, V, C):  
2.   sort E by C  
3.   T = {}  
4.   for i = 1 to m:  
5.     if T U {ei} has no cycle:  
6.       T = T U {ei}
```

- **Assumption:**
 - G is connected
 - Distinct c_e values

Example



Proof of Correctness

- **Claim:** Kruskal's algorithm correctly computes an MST
- **Part I:** Kruskal's algorithm produces a spanning tree T^*
 - Kruskal's creates graph that spans G . (for $i = 1$ to m :)
 - Are they tree? (no cycle)
 - Is it connected?
- **Part III:** T^* is a MST
 - Is it minimal?

Kruskal's is a spanning tree

- **Part Ia:** T^* has no cycle
 - By the nature of the algorithm, T^* does not create a cycle.
- **Part Ib:** T^* is connected (single tree)
 - To prove that T^* is connected, we need to prove that every cut of T^* has at least one edge that crosses it. (By Empty Cut Lemma)
 - Consider a cut of T^* , since G is connected there must be at least one edge that crosses the cut that's part of G .
 - Of those edge, the cheapest edge would be part of T^* because that's the first edge of that cut.

Kruskal's is Minimal

- Every edge of T^* is justified by the Cut Property
- **The Cut Property**
 - Given $e \in G$, suppose \exists cut (A,B) | e is the cheapest crossing edge, then $e \in \text{MST}(G)$

Proof of Part II


- At each step of the algorithm, we have a forest.
- Let (u,v) be the edge that the algorithm picks.
- Since $T \cup \{ (u,v) \}$ has no cycle $\Rightarrow \exists$ an empty cut, (A, B) , that separates u and v .
- (u,v) must be the cheapest edge that crosses (A,B) because all other edge that was not part of the forest cost more than (u,v) .

Running time analysis

- Naive implementation
 - Checking for cycle = $O(n)$ using DFS or BFS.
 - $O(m \log m) + O(mn) = O(mn)$
pre-processing cycle checks total running time
- Can we do better?
 - Speeding cycle checking by using a Union-Find data structure.

Union-Find

- **Purpose:** maintain partition of a set of objects
- **Operations:**
 - $\text{Find}(x)$ = returns the group that x belongs to
 - $\text{Union}(C_i, C_j)$ = combine C_i and C_j into single group
 - Name of a group can be represented by its representative.
- Kruskal's Application
 - Objects = vertices
 - Groups = connected components (Forests)



Kruskal with Union-Find

```
1. Kruskal(V, E, C):
2.   Sort E by C
3.    $\forall v \in V$  is in its own group
4.    $T = \{\}$ 
5.   for  $e_i = e_1$  to  $e_m$ :
6.      $(u, v) = e_i$ 
7.     if FIND(u)  $\neq$  FIND(v):
8.        $T = T \cup \{e_i\}$ 
9.       UNION(u, v)
10.  return T
```

- Running Time

- $O(m \log m) + O(m (\text{FIND}(n) + \text{UNION}(n)))$

Naive Union-Find implementation

```
1. UNION(a, b):  
2.     FIND(b).parent = a
```

```
1. FIND(b):  
2.     if b.parent == b:  
3.         b  
4.     else  
5.         FIND(b.parent)
```

- Running time?

- UNION: $O(n)$

- FIND: $O(n)$

- Kruskal's Running Time

- $O(m \log m) + O(m (O(1) + O(n)))$

- $= O(mn)$

Eager Union

```
1. UNION(a,b):  
2.     if SIZE(a) > SIZE(b):  
3.         UNION(b,a)  
4.     for i in SET(a):  
5.         i.parent = b  
  
7. FIND(x):  
8.     x.parent
```

- $\text{UNION}(a,b)$
 - Take the smaller set and make all its element point to the root of the other set
 - $O(n)$??
- $\text{FIND}(x)$
 - $O(1)$

Eager Union Analysis

- Single UNION call is bounded by $O(n)$. However, in our usage in Kruskal's algorithm, we end up calling UNION until the entire vertices are merged into one set.
- What is the amortized (entire execution) running time of UNION then?
- Fix an element, x . We want to know how many times we have to update its parents during the course of all UNIONS. Whenever x has to update its parent, the target set is always at least double the size.
- So finally, the question is "how many times can a number double before it is the size of n ?"
- $\text{UNION}(a, b) = O(\log n)$

Lazy Union

```
1.  UNION(a,b):
2.    x = FIND(a)
3.    y = FIND(b)
4.    if RANK(x) > RANK(y):
5.        y.parent = x
6.    else:
7.        x.parent = y
8.    update RANK

10. FIND(x):
11.   if (x.parent == x):
12.       return x
13.   else:
14.       return FIND(x.parent)
```

- Take the root of the smaller tree and make it the child of the other
- Let $\text{RANK}[x] = \text{maximum \# of hops from leaf to } x$
- Initially $\text{RANK}[x] = 0$ for all possible x
- Analysis
 - $\text{UNION} = O(1)$
 - $\text{FIND} = O(n)??$
 - Note, $\text{max RANK} = \text{worst case running time for FIND}$

Lazy Union Analysis

- Note:
 - For all object x , $\text{RANK}[x]$ only goes up over time
 - Only ranks of the root can go up
 - $\text{RANK}[x]$ strictly increases along the path to the root

Rank Lemma

- Consider an arbitrary sequence of UNIONS
- $\forall r \in \{0, 1, 2, \dots\}$ there are at most $n/2^r$ object with rank r
- Corollary:
 - if $r = \log_2 n \rightarrow n/2^{\log_2 n} = 1$
 - $\max \text{rank} \leq \log_2 n$
 - $\text{FIND, UNION} = O(\log n)$
 - Kruska's Running Time: $O(m \log m) + O(m (\log n + \log n))$
 $= O(m \log n)$

Proof of Rank Lemma

- Claim 1
 - If x, y have same rank, then their subtree are disjoint
- Claim 2
 - The subtree of a rank r has object size $\geq 2^r$
- Fix an r . Each of the object with rank r has 2^r objects that can reach them. Since they are all disjoint and there are at most n objects number of objects with the r must be bounded by $n/2^r$

Proof of Rank Lemma

- Claim 1
 - If x, y have same rank, then their subtree are disjoint
- Proof by Contradiction
 - If x, y are not disjoint, $\exists z \mid \exists \text{ paths } z \rightsquigarrow x, z \rightsquigarrow y$
 - Because of how trees are structured, the path must be either
$$z \rightsquigarrow x \rightsquigarrow y \text{ or } z \rightsquigarrow y \rightsquigarrow x$$
 - Since rank increases strictly, this is a contradiction
 - QED

Proof of Rank Lemma

- Claim 2
 - The subtree of a rank r has object size $\geq 2^r$
- Proof by Induction
 - Base Case: rank = 0, all subtree size = 1
 - Inductive Hypothesis: if $\text{SIZE}[k] \geq 2^k$ then $\text{SIZE}[k+1] \geq 2^k$
 - Inductive Case: $\text{UNION}(a, b) \rightarrow x = \text{FIND}(a), y = \text{FIND}(b)$
 - $\text{RANK}[a] = \text{RANK}[b] = r$
 - b 's new rank = $r + 1$
 - b 's new size = $\text{SIZE}[b] + \text{SIZE}[a] \geq 2^r + 2^r = 2 \cdot 2^r = 2^{r+1}$

Can we do better?

- Path Compression
 - We do awful lot of FIND
 - So whenever we call $\text{FIND}[x]$, we update all the object along the path from x to root such that their parents are now pointing to the root.
- Analysis
 - $O(\log^* n)$ where $\log^* n := \begin{cases} 0 & \text{if } n \leq 1; \\ 1 + \log^*(\log n) & \text{if } n > 1 \end{cases}$
 - $\log^* (2^{65,536}) = 5$
 - For all $n < 65,536$, $\text{Kruskal} = O(m \log n) + O(m (5 + 5))$

Can we do better?

- Randomized UNION-FIND = $O(m)$
- Deterministic algorithm that runs in $O(m)$?
- Best so far = $O(m \alpha(n))$ where α is an inverse Ackermann function
 - It grows slower than \log^*
- 2002 Paper
 - There exists an algorithm that has been proven to be the optimal algorithm. But its running time hasn't been proven.
 - All we know is that it runs between $\Theta(m)$ and $O(m \alpha(n))$
- Open question
 - Simple randomized UNION-FIND?
 - Deterministic $O(m)$?